

# From Exceptions to Pull Requests: Building a Self Fixing .NET 10 App with AI Agents

\*\*\*\*



## What if your application could fix itself?

Not science fiction. Not a distant future. Today, with AI agents, we can build systems that:

- Monitor Application Insights 24/7
- Analyze exceptions and read your actual source code
- Diagnose root causes with LLM reasoning
- Propose fixes with clear explanations
- Create Pull Requests automatically

I built one. Here I used Azure services but you can do the same by using others.

In this article, we'll cover :

- The Infrastructure to manage a self-fixing production application
- How it works
- The cost

**Try it yourself here** (includes Bicep templates files to create new Azure Resources):

<https://github.com/jBeyondstars/ai-diagnostics-agent>

Every time an exception appears in my application's App Insights (Azure's service for logs \*\*treatment, management, ...\*\* my Agent API is called to analyze and process the issue.

## What Makes an AI Agent Different?

Before we dive into code, let's clarify an important distinction: an AI agent is not a chatbot.

A chatbot answers questions. An agent takes actions.

Think of it this way:

Chatbot: "Here's how you might fix that NullReferenceException..."

Agent: reads your code, identifies the bug, creates a branch, commits a fix, opens a PR

The agent follows a loop: Perceive → Reason → Act → Observe.

It perceives the world through its tools. (\*\*example of tools\*\*)

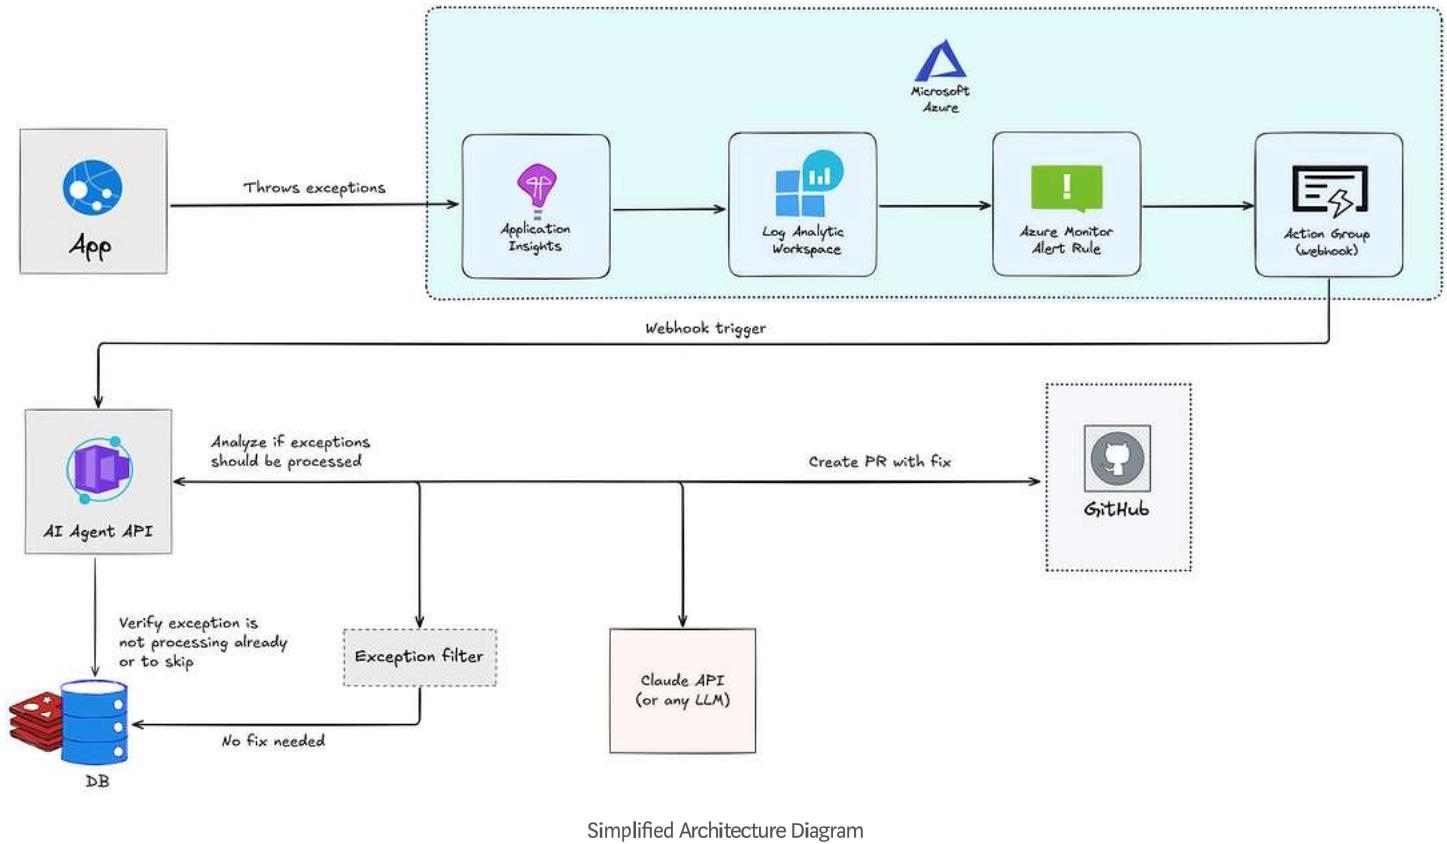
We will use Semantic Kernel (and explain what it is) to define "plugins" for the collections of functions the agent can call. The LLM decides which functions to call and in what order.

We don't write IF/ELSE logic; the agent will reason.

Our agent has two "senses":

- AppInsights Plugin: Sees exceptions, stack traces, -performance
- Metrics GitHub Plugin: Reads code, creates branches, opens PRs

## The Architecture



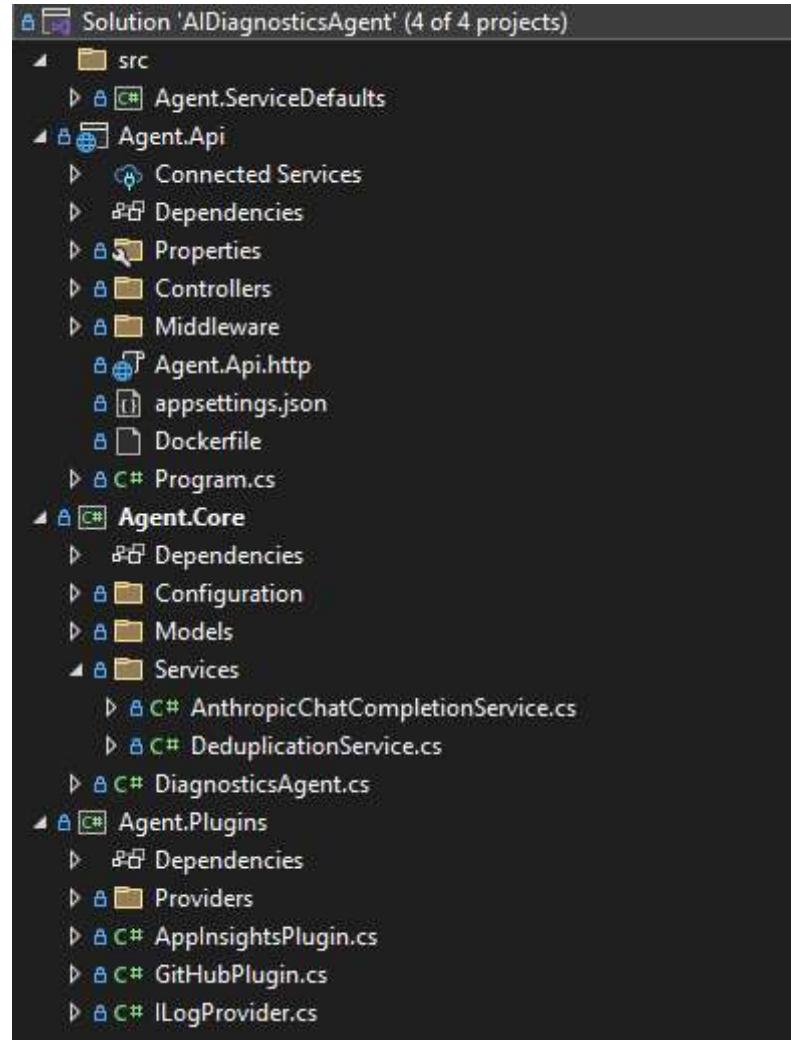
As I said, for this project I used an Azure environment but you could apply the same process in your own setup.

Let's deep dive into the infrastructure :

- A .NET 10 App that is being monitored
- Azure Monitor / AppInsights / Log Analytics Workspace / Alert Rule & Action Group
- A .NET 10 Agent App
- A Redis In memory cache but I highly recommend the use of an external DB
- Claude API (we could go with Azure OpenAI SDK for fully Azure integrated or any modern LLM API such as Gemini or ChatGPT). I personally prefer Claude Opus 4.5 model for coding task.
- A GitHub Repo and Personal Access Token with permissions to allow our agent to create the PRs.

For the demo, I used a single solution for both the App and the agent.

It is structured as below:



## Agent.Api

The REST API layer that exposes our diagnostic endpoints.

When Azure Monitor detects exceptions, it calls the `/analyze` webhook endpoint here.

## Agent.Core

The brain of the operation. This is where the orchestration happens, among other things:

- Creates and manages the Semantic Kernel agent
- Prevents analyzing the same exception twice

## Agent.Plugins

The custom tools that Claude can invoke. I set up a list of tools for communicating with Azure AppInsights and GitHub. Find the list below.

AppInsights Plugin:

Function Name	Description
get_exceptions	Retrieves exceptions grouped by type with occurrence counts
get_exception_details	Gets detailed samples with full stack traces

GitHub Plugin:

Tool	Purpose
get_file_content	Retrieves source file content from the repository
search_code	Searches for code snippets in the repository
get_multiple_files	Retrieves multiple files efficiently in one call

....

## The Flow

Let's take an example to see what it exactly does.

Mr. Doe is a user of the app. He ... \*\*find scenario\*\*

A few seconds later, App Insights will show the exceptions created by the App when the errors occurred. App Insights is the place you can interact with logs and exception that are ingested from the Log Analytic Workspace (LAW)

The screenshot shows the 'End-to-end transaction details' view in the Azure App Insights portal. It highlights an exception event within a transaction. The exception is identified as 'System.DivideByZeroException'. The call stack provides detailed information about the error's origin, including multiple layers of Microsoft .NET Core framework code.

App Insights interface on exception details

Aside, our Azure Monitor Alert Rule is configured to query the LAW using KQL (Kusto Query Language) every 5 minutes.

The query we are using is pretty simple:

```
AppExceptions
| where TimeGenerated > ago(5m)
| summarize Count = count() by ExceptionType, Pr
| where Count >= 1
```

Ding, after a few seconds when the query is triggered, the alert \*\*verb\*\*.

We can verify on the Azure Portal in our Alert Rule history:

Severity	Fired time	Affected resource	Monitor service	Alert condition	User response
2 - Warning	1/2026, 8:18 PM	log-zgqx4valgxhdu	Log Alerts V2	Fired	New

This Alert Rule has one job. Calling the /analyze endpoint of our Agent API (webhook).

The endpoint will run the analyze of the exceptions that occurred the last 5 minutes. If many occurred, it will handle the remediation of all

of them separately.

If an exception is already being treated or as been marked as skipped (Exception filter) in the DB, we just pass.

The Exception Filter now takes place. Its role, as \*\*en anglais : comme son nom le fait penser\*\* is to filter the exception that should not be treated. We will deep dive in that in the next section.

Once we know what exception to fix, we prepare a fully detailed prompt for the LLM:

```
"You are an expert DevOps engineer and .NET deve  
production exceptions from Application Insights
```

### ## Your Capabilities

You have access to these tools:

- **AppInsights**: Query exceptions, failed requ
- **Github**: Read source code, search for files

### ## Your Workflow

1. **Discover**: Use `get\_exceptions` to find th
2. **Investigate**: For each critical exception:
  - Analyze the stack trace to identify the sou
  - Use `get\_file\_content` to read the relevant
  - Use `search\_code` if you need to find relat
3. **Diagnose**: Determine the root cause of eac
4. **Fix**: Generate corrected code that fixes t
5. **Document**: Create a PR with clear explanat

### ## Guidelines

- Focus on exceptions with HIGH occurrence count
- Always read the actual source code before prop
- Consider null checks, exception handling, and
- Provide clear explanations of what was wrong a
- If you're not confident about a fix, create an
- Never make changes that could break existing f
- Follow C# best practices and the existing code

### ## Output Format

When generating fixes, always explain:

1. What exception was occurring and why
2. What the root cause is in the code

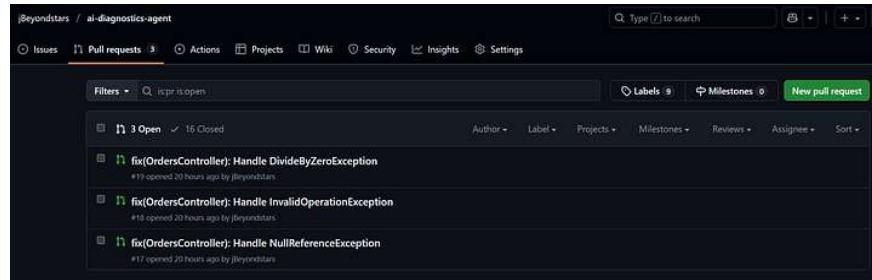
3. How your fix addresses the root cause
4. Any tests that should be added

Be thorough but concise. Focus on actionable fix

Here we explain to Claude what he tools he can use and how we want him to proceed with the exceptions.

We'll deep dive in the Tools section to understand how Claude understands how to use the tools.

The remediation is now prepared and ready to ship. The PRs (one for each exception) are created by our agent using the plugins we created.



The Agent created a PR for each fixable exception

We can detail the PR description and title as much as we want. Our agent currently provide a brief summary, some exception details with a link to the exception on App Insights.

Fully detailed description with App Insights link of the exception

You may wonder know, is the fix viable? Well, if you ever worked with Claude (especially Opus 4.5) you already know how reliable it is in code generation.

So yes, every fixes suggested so far in my tests were consistent and pertinent.

The changes suggested by the Agent for a piece of code trying to call `.First()` on a collection without verifying in contains elements.

## Exception Filtering

As you know, all exceptions don't require a remediation.  
`HttpRequestException` when an external API is down? No code fix.  
`TimeoutException` on a slow database? Infrastructure problem.

The agent uses a 3-level filtering system.

Level 1: Exclusion List Infrastructure and transient exceptions are automatically skipped

```
public string[] ExcludedTypes { get; set; } =
[
    "System.Net.Http.HttpRequestException",
    "System.TimeoutException",
    "System.Threading.Tasks.TaskCanceledException",
    "StackExchange.Redis.RedisConnectionException",
    "Azure.RequestFailedException",
    "Microsoft.Data.SqlClient.SqlException",
    // ...
];
```

Level 2: Pattern Matching Partial string matches catch variants

```
public string[] ExcludedPatterns { get; set; } =
[
    "ConnectionException",
    "Timeout",
    "Unreachable",
    "NetworkError",
    "HostNotFound"
];
```

Level 3: Claude Evaluation

For ambiguous exceptions (containing “api”, “service”, “500”), Claude evaluates.

```
private async Task<bool> EvaluateIfFixableAsync(
{
    var prompt = $$"
        Exception Type: {exception.ExceptionType}
        Message: {exception.Message}
    $$";
}
```

```
        Can this exception be prevented by modif
        Consider:
        - Transient infrastructure error → NO
        - Missing error handling or validation →
        - Bug in business logic → YES
        - External service down → NO
        Respond with ONLY: YES or NO
        """"";
    }
```

...

## How Tool Use Works with Claude API

When building an AI agent that can interact with external systems, you might think that simply telling Claude about available tools in the system prompt is enough. It's not.

Here's what actually happens behind the scenes.

### 1. Tool Definitions

When calling the Claude API, we must pass a tools array containing JSON schemas for each tool. This tells Claude what tools exist and how to call them:

```
{
  "name": "AppInsights_get_exceptions",
  "description": "Retrieves exceptions from th
  "input_schema": {
    "type": "object",
    "properties": {
      "hours": { "type": "integer", "default": 24 }
    }
  }
}
```

### 2. Tool Execution

Our backend must contain the actual code that performs the work.  
In our case, we use Semantic Kernel plugins:

```
[KernelFunction("get_exceptions")]
[Description("Retrieves exceptions from the la
public async Task<string> GetExceptionsAsync(i
{
    // Actually query Application Insights her
    // Code available in the Github repo
}
```

## What is Kernel?

Semantic Kernel is Microsoft's SDK that acts as a bridge between your code and LLMs. Think of it as the middleware that connects Claude (or any LLM) to your application's capabilities.

In this project we use the Kernel for three purposes:

- **LLM Abstraction:** The Kernel lets us swap between Claude, OpenAI, or Azure OpenAI with a simple config change.
- **Plugin Registry**
- **Agent Orchestration:** The Kernel is passed to the Agent, which uses it to discover and execute tools

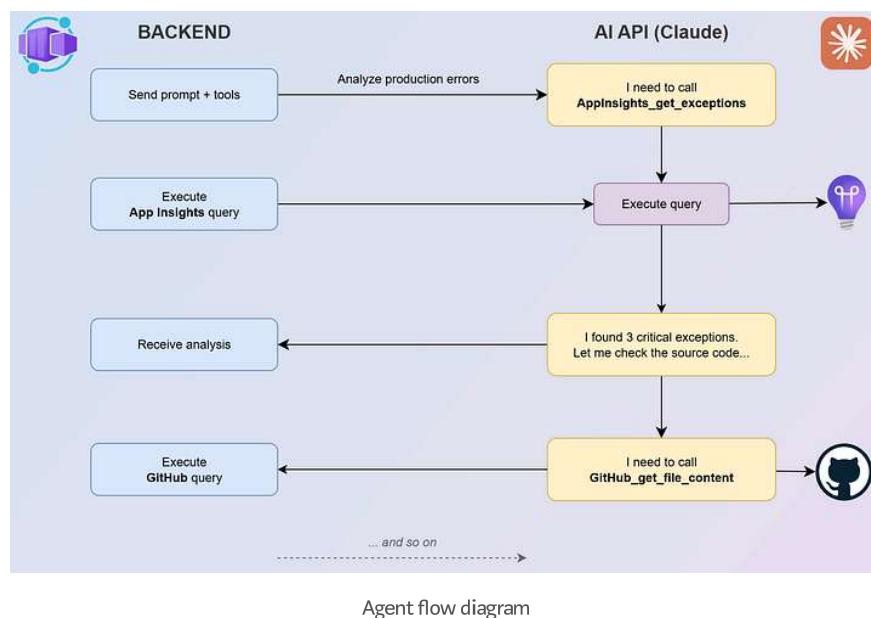
```
var agent = new ChatCompletionAgent
{
    Name = "DiagnosticsAgent",
    Instructions = SystemPrompt,
    Kernel = kernel, // Contains all registered p
    Arguments = new KernelArguments(new OpenAIPro
    {
        FunctionChoiceBehavior = FunctionChoiceBehavi
    })
};
```

The `FunctionChoiceBehavior.Auto()` is the magic. The LLM autonomously decides which tools to invoke based on the conversation context.

### 3. The Tool Loop

When Claude decides to use a tool, it doesn't execute anything. It simply requests a tool call. Our code must:

1. Detect that Claude wants to use a tool (`StopReason = "tool_use"`)
2. Parse the tool name and arguments
3. Execute the corresponding function
4. Send the result back to Claude
5. Let Claude continue reasoning with the new data



## Deduplication Handling

I used Redis for this project but my thoughts is the best would be an hybrid model with Redis and a Database.

....

....

As we store the ongoing fixes in our side, we need to configure GitHub to know if a PR is manually closed (merged or canceled). For that we use a simple webhook on action and call an endpoint of our agent to update in our side.

## Cost Breakdown

I was actually agreeably surprised by how cheap the whole process is.

### Per-Analysis Cost

Claude API cost

Alert rule cost:

#### Alert logic

Threshold type  Static  Dynamic

Operator  Greater than or equal to

Threshold value \*

Frequency of evaluation \*

 Estimated monthly cost **\$1.50 (USD)**

...

### ROI Calculation

....

## What's Next?

The Agent is a prototype still, there are few enhancements that could and should be added to be production ready.

1. Multi-Agent (Coder Agent, Reviewer Agent..)
2. Rework on Pull Request comments (Agent would listen to your comments and suggest new fix)
3. Slack/Teams Integration to notify when PRs are created Allow manual approval before creation

4. Learning Loop Track which fixes were merged vs. rejected Fine-tune the prompt based on feedback
- 5.

## Conclusion

We built a complete self-healing system that:

- Monitors 24/7 via Azure Monitor
- Alert Rules Triggers automatically via webhooks
- - Filters non-actionable exceptions (3-level smart filtering)
- Deduplicates at multiple levels (Redis + existing PR check)
- Analyzes with Claude in a single LLM call
- Creates Pull Requests automatically
- Costs less than ~\$0.01 per analysis

## Resources

.